AI COMPUTE

Total points 58/60

60 QUESTIONS. ALL QUESTIONS ARE COMPULSORY

Email *

patsonu63@gmail.com

NAME & PRN

Sumit Suresh patil 220980728007

COURSE

- [ ] DAC
- [x] DAI
- [ ] DBDA



CENTRE

- ○ 1. BENGALURU
- ○ 2. CHENNAI
- ○ 3. GUWAHATI
- ○ 4. HYDERABAD
- ○ 5. KARAD
- ○ 6. KOCHI
- ○ 7. KOLKATA
- ○ 8. MUMBAI
- ○ 9. NASHIK
- ○ 10. NOIDA
- ◉ 11. PATNA
- ○ 12. PUNE
- ○ 13. SILCHAR
- ○ 14. THIRUVANANTHAPURAM



✓ 1. Spark was initially started by _____ at UC Berkeley AMPLab in 2009. 1

- a) Mahek Zaharia
- b) Matei Zaharia
- c) Doug Cutting
- d) Stonebraker

✓ 2. Point out the correct statement. * 1

- a) RSS abstraction provides distributed task dispatching, scheduling, and basic I/ functionalities
- b) For cluster manager, Spark supports standalone Hadoop YARN
- c) Hive SQL is a component on top of Spark Core
- d) None of the mentioned

✓ 3. _____ is a component on top of Spark Core. * 1

- a) Spark Streaming
- b) Spark SQL
- c) RDDs
- d) All of the mentioned

✓ 4. Spark SQL provides a domain-specific language to manipulate _____ in Scala, Java, or Python. 1

- a) Spark Streaming
- b) Spark SQL
- (●) c) RDDs
- d) All of the mentioned

✓ 5. Point out the wrong statement. * 1

- a) For distributed storage, Spark can interface with a wide variety, including Hado Distributed File System (HDFS)
- b) Spark also supports a pseudo-distributed mode, usually used only for development or testing purposes
- c) Spark has over 465 contributors in 2014
- (●) d) All of the mentioned

✓ 6. _____ leverages Spark Core fast scheduling capability to perform streaming analytics. 1

- a) MLlib
- (●) b) Spark Streaming
- c) GraphX
- d) RDDs

✓ 7. _____ is a distributed machine learning framework on top of Spark. 1

- (●) a) MLlib
- ( ) b) Spark Streaming
- ( ) c) GraphX
- ( ) d) RDDs

✓ 8. _____ is a distributed graph processing framework on top of Spark. * 1

- ( ) a) MLlib
- ( ) b) Spark Streaming
- (●) c) GraphX
- ( ) d) All of the mentioned

✓ 9. GraphX provides an API for expressing graph computation that can model the _____ abstraction. 1

- ( ) a) GaAdt
- ( ) b) Spark Core
- (●) c) Pregel
- ( ) d) None of the mentioned

✓ 10. Spark architecture is _____ times as fast as Hadoop disk-based Apache Mahout
and even scales better than Vowpal Wabbit. 1/1

- (●) a) 10
- ( ) b) 20
- ( ) c) 50
- ( ) d) 100

✓ 11. Users can easily run Spark on top of Amazon's _____ * 1/1

- ( ) a) Infosphere
- (●) b) EC2
- ( ) c) EMR
- ( ) d) None of the mentioned

✓ 12. Point out the correct statement. * 1/1

- (●) a) Spark enables Apache Hive users to run their unmodified queries much faster
- ( ) b) Spark interoperates only with Hadoop
- ( ) c) Spark is a popular data warehouse solution running on top of Hadoop
- ( ) d) None of the mentioned

✓ 13. Spark runs on top of _____ a cluster manager system which provides efficient
resource isolation across distributed applications. 1

- a) Mesjs
- b) Mesos
- c) Mesus
- d) All of the mentioned

✓ 14. Which of the following can be used to launch Spark jobs inside MapReduce? 1

- a) SIM
- b) SIMR
- c) SIR
- d) RIS

✓ 15. Point out the wrong statement. * 1

- a) Spark is intended to replace, the Hadoop stack
- b) Spark was designed to read and write data from and to HDFS, as well as other storage systems
- c) Hadoop users who have already deployed or are planning to deploy Hadoop Y can simply run Spark on YARN
- d) None of the mentioned

✓ 16. Which of the following language is not supported by Spark? * 1

- a) Java
- b) Pascal
- c) Scala
- d) Python

✓ 17. Spark is packaged with higher level libraries, including support for _____ queries. 1

- a) SQL
- b) C
- c) C++
- d) None of the mentioned

✓ 18. Spark includes a collection over _____ operators for transforming data and familiar
data frame APIs for manipulating semi-structured data. 1

- a) 50
- b) 60
- c) 70
- d) 80

✓ 19. Spark is engineered from the bottom-up for performance, running _____ faster than Hadoop by exploiting in memory computing and other optimizations. 1/1

- (●) a) 100x
- ( ) b) 150x
- ( ) c) 200x
- ( ) d) None of the mentioned

✓ 20. Spark powers a stack of high-level tools including Spark SQL, MLlib for _____ 1/1

- ( ) a) regression models
- ( ) b) statistics
- (●) c) machine learning
- ( ) d) reproductive research

✓ 21. Apache Flume 1.3.0 is the fourth release under the auspices of Apache of the so-called _____ codeline. 1

- (●) a) NG
- ( ) b) ND
- ( ) c) NF
- ( ) d) NR

✓ 22. Point out the correct statement. 1

- (●) a) Flume is a distributed, reliable, and available service
- ( ) b) Version 1.5.2 is the eighth Flume release as an Apache top-level project
- ( ) c) Flume 1.5.2 is production-ready software for integration with hadoop
- ( ) d) All of the mentioned

✓ 23. _____ was created to allow you to flow data from a source into your Hadoop environment. 1

- ( ) a) Imphala
- ( ) b) Oozie
- (●) c) Flume
- ( ) d) All of the mentioned

✓ 24. A _____ is an operation on the stream that can transform the stream. 1

- a) Decorator
- (•) b) Source
- c) Sinks
- d) All of the mentioned

✓ 25. Point out the wrong statement. 1

- a) Version 1.4.0 is the fourth Flume release as an Apache top-level project
- b) Apache Flume 1.5.2 is a security and maintenance release that disables SSLv3 on all components in Flume that support SSL/TLS
- c) Flume is backwards-compatible with previous versions of the Flume 1.x codel
- (•) d) None of the mentioned

✓ 26. A number of _____ source adapters give you the granular control to grab a specific file. 1

- a) multimedia file
- (•) b) text file
- c) image file
- d) none of the mentioned

✓ 27. _____ is used when you want the sink to be the input source for another operation. 1

- a) Collector Tier Event
- b) Agent Tier Event
- c) Basic
- d) All of the mentioned

✓ 28. _____ is where you would land a flow (or possibly multiple flows joined together) into an HDFS-formatted file system. 1

- a) Collector Tier Event
- b) Agent Tier Event
- c) Basic
- d) All of the mentioned

✓ 29. _____ sink can be a text file, the console display, a simple HDFS path, or a null bucket where the data is simply deleted. 1

- a) Collector Tier Event
- b) Agent Tier Event
- c) Basic
- d) None of the mentioned

✓ 30. Flume deploys as one or more agents, each contained within its own instance of_____ 1

- a) JVM
- b) Channels
- c) Chunks
- d) None of the mentioned

✓ 31. _____ provides Java-based indexing and search technology. * 1

- a) Solr
- b) Lucene Core
- c) Lucy
- d) All of the mentioned

✓ 32. Point out the correct statement. * 1

- a) Building PyLucene requires GNU Make, a recent version of Ant capable of building Java Lucene and a C++ compiler
- b) PyLucene is supported on Mac OS X, Linux, Solaris and Windows
- c) Use of setuptools is recommended for Lucene
- (•) d) All of the mentioned

✓ 33. _____ is a high performance search server built using Lucene Core. 1

- (•) a) Solr
- b) Lucene Core
- c) Lucy
- d) PyLucene

✓ 34. _____ is a subproject with the aim of collecting and distributing free materials. 1

- a) OSR
- b) OPR
- (•) c) ORP
- d) ORS

✓ 35. Point out the wrong statement. * /1

- (●) a) PyLucene is a Lucene port
- ( ) b) PyLucene embeds a Java VM with Lucene into a Python process
- ( ) c) The PyLucene Python extension, a Python module called lucene is machine-generated by JCC
- ( ) d) PyLucene is built with JCC

✓ 36. _____ is a Python port of the Core project. * /1

- ( ) a) Solr
- ( ) b) Lucene Core
- ( ) c) Lucy
- (●) d) PyLucene

✓ 37. The Lucene _____ is pleased to announce the availability of Apache Lucene 5.0.0 and Apache Solr 5.0.0. /1

- (●) a) PMC
- ( ) b) RPC
- ( ) c) CPM
- ( ) d) All of the mentioned

✓ 38. _____ is a technology suitable for nearly any application that requires full-text search, especially cross-platform. 1

- (●) a) Lucene
- ( ) b) Oozie
- ( ) c) Lucy
- ( ) d) All of the mentioned

✓ 39. Lucene provides scalable, high-Performance indexing over _____ per hour on modern hardware. 1

- ( ) a) 1 TB
- (●) b) 150GB
- ( ) c) 10 GB
- ( ) d) None of the mentioned

✓ 40. Lucene index size is roughly _____ the size of text indexed. * 1

- ( ) a) 10%
- (●) b) 20%
- ( ) c) 50%
- ( ) d) 70%

✓ 41. _____ is the architectural center of Hadoop that allows multiple data processing engines. 1

- (●) a) YARN
- ( ) b) Hive
- ( ) c) Incubator
- ( ) d) Chuckwa

✓ 42. Point out the correct statement. * 1

- ( ) a) YARN also extends the power of Hadoop to incumbent and new technologies found within the data center
- ( ) b) YARN is the central point of investment for Hortonworks within the Apache community
- ( ) c) YARN enhances a Hadoop compute cluster in many ways
- (●) d) All of the mentioned

✓ 43. YARN's dynamic allocation of cluster resources improves utilization over more static _____ rules used in early versions of Hadoop. 1

- ( ) a) Hive
- (●) b) MapReduce
- ( ) c) Imphala
- ( ) d) All of the mentioned

✓ 44. The _____ is a framework-specific entity that negotiates resources from the Resource Manager. 1

- a) NodeManager
- b) ResourceManager
- (●) c) ApplicationMaster
- d) All of the mentioned

✓ 45. Point out the wrong statement. * 1

- a) From the system perspective, the ApplicationMaster runs as a normal contain
- (●) b) The ResourceManager is the per-machine slave, which is responsible for launching the applications' containers
- c) The NodeManager is the per-machine slave, which is responsible for launchin the applications' containers, monitoring their resource usage
- d) None of the mentioned

✓ 46. Apache Hadoop YARN stands for _____ * 1

- a) Yet Another Reserve Negotiator
- b) Yet Another Resource Network
- (●) c) Yet Another Resource Negotiator
- d) All of the mentioned

✓ 47. MapReduce has undergone a complete overhaul in hadoop is _____ 1

- a) 0.21
- b) 0.23
- c) 0.24
- d) 0.26

✓ 48. The _____ is the ultimate authority that arbitrates resources among all the applications in the system. 1

- a) NodeManager
- b) ResourceManager
- c) ApplicationMaster
- d) All of the mentioned

✓ 49. The _____ is responsible for allocating resources to the various running applications subject to familiar constraints of capacities, queues etc. 1

- a) Manager
- b) Master
- c) Scheduler
- d) None of the mentioned

✓ 50. The CapacityScheduler supports _____ queues to allow for more predictable
sharing of cluster resources. 1

- a) Networked
- b) Hierarchical
- c) Partition
- d) None of the mentioned

✓ 51. Yarn commands are invoked by the _____ script. * 1

- a) hive
- b) bin
- c) hadoop
- d) home

✓ 52. Point out the correct statement. 1

- a) Each queue has strict ACLs which controls which users can submit applicatio to individual queues individual queues
- b) Hierarchy of queues is supported to ensure resources are shared among the s queues of an organization
- c) Queues are allocated a fraction of the capacity of the grid in the sense that a certain capacity of resources will be at their disposal
- d) All of the mentioned

✗ 53. The queue definitions and properties such as _____ ACLs can be changed, at runtime. 1

- a) tolerant
- b) capacity
- c) speed
- d) all of the mentioned

Correct answer

- b) capacity

✓ 54. The CapacityScheduler has a predefined queue called _____ * 1

- a) domain
- b) root
- c) rear
- d) all of the mentioned

✗ 55. Point out the wrong statement. 1

- a) The multiple of the queue capacity which can be configured to allow a single user to acquire more resources
- b) Changing queue properties and adding new queues is very simple
- c) Queues cannot be deleted, only addition of new queues is supported
- d) None of the mentioned

Correct answer

- d) None of the mentioned

✓ 56. The updated queue configuration should be a valid one i.e. queue-capacity at each level
should be equal to _____ 1

- a) 50%
- b) 75%
- c) 100%
- d) 0%

✓ 57. Users can bundle their Yarn code in a _____ file and execute it using jar command. 1

- a) java
- b) jar
- c) C code
- d) xml

✓ 58. Which of the following command is used to dump the log container? * 1

- a) logs
- b) log
- c) dump
- d) all of the mentioned

✓ 59. _____ will clear the RMStateStore and is useful if past applications are no longer needed. 1

- a) -format-state
- b) -form-state-store
- c) -format-state-store
- d) none of the mentioned

✓ 60. Which of the following command runs ResourceManager admin client? 1

- a) proxyserver
- b) run
- c) admin
- d) rmadmin

This content is neither created nor endorsed by Google. - Terms of Service - Privacy Policy

Google Forms

